

DATA 442: Neural Networks & Deep Learning

Dan Runfola – danr@wm.edu
icss.wm.edu/data442/



34	40	34	3	8	30	50	26	16	28	41	6	23	2	24	0
14	37	19	25	6	9	14	17	4	46	20	7	38	6	29	28
4	39	0	29	15	6	50	2	21	10	8	45	150	145	106	46
42	10	15	19	24	18	111	123	118	104	119	122	117	140	138	28
21	35	19	30	14	143	146	147	142	103	109	127	108	148	20	23
30	105	147	102	126	118	108	101	140	131	124	136	47	27	26	38
135	133	137	108	140	144	135	120	118	137	125	43	8	31	45	10
106	142	108	138	137	111	38	36	32	1	19	44	34	4	38	49
122	142	127	131	143	8	47	4	0	31	39	18	46	1	50	25
149	137	122	36	50	19	24	45	16	30	2	47	2	35	29	50
147	115	3	29	10	2	13	1	48	3	45	28	39	14	14	20





CAOE | CENTER FOR ACCELERATING OPERATIONAL EFFICIENCY

A DEPARTMENT OF HOMELAND SECURITY CENTER OF EXCELLENCE

34	40	34	3	8	30	50	26	16	28	41	6	23	2	24	0
14	37	19	25	6	9	14	17	4	46	20	7	38	6	29	28
4	39	0	29	15	6	50	2	21	10	8	45	150	145	106	46
42	10	15	19	24	18	111	123	118	104	119	122	117	140	138	28
21	35	19	30	14	143	146	147	142	103	109	127	108	148	20	23
30	105	147	102	126	118	108	101	140	131	124	136	47	27	26	38
135	133	137	108	140	144	135	120	118	137	125	43	8	31	45	10
106	142	108	138	137	111	38	36	32	1	34	40	34	3	8	30
122	142	127	131	143	8	47	4	0	31	14	37	19	25	6	9
149	137	122	36	50	19	24	45	16	30	28	6	31	39	0	23
147	115	3	29	10	2	13	1	48	3	71	6	9	44	41	23

34	40	34	3	8	30	50	26	16	28	41	6	23	2	24	0
14	37	19	25	6	9	14	17	4	46	20	7	38	6	29	28
28	6	31	39	0	23	36	34	21	10	8	45	6	29	45	46
71	6	9	44	41	23	36	61	4	104	119	122	117	140	138	28
8	60	45	11	12	165	122	115	142	103	109	127	108	148	100	23
2	94	156	88	174	160	62	59	140	131	124	136	127	150	145	38
183	94	160	101	108	163	135	119	118	137	125	43	8	31	45	10
128	179	74	122	89	140	59	22	32	1	19	44	34	4	38	49
165	100	106	172	110	41	58	11	0	31	39	18	46	1	50	25
154	175	158	76	53	58	61	69	16	30	2	47	2	35	29	50
153	138	43	12	19	40	62	26	48	3	45	28	39	14	14	20



CAOEI | CENTER FOR
OPERATIONAL
A DEPARTMENT OF HOMELAND SECURITY



**Intra-Class
Differences**

CAOE | CENTER FOR ACCELERATING
OPERATIONAL EFFICIENCY

A DEPARTMENT OF HOMELAND SECURITY CENTER OF EXCELLENCE



Viewpoint



**Intra-Class
Differences**



CAOE | CENTER FOR ACCELERATING
OPERATIONAL EFFICIENCY


A DEPARTMENT OF HOMELAND SECURITY CENTER OF EXCELLENCE



Viewpoint



**Intra-Class
Differences**



Background

CENTER FOR ACCELERATING
 OPERATIONAL EFFICIENCY
 OF HOMELAND SECURITY CENTER OF EXCELLENCE




Viewpoint



Lighting



Intra-Class Differences



Background

CENTER FOR ACCELERATING
 OPERATIONAL EFFICIENCY
 OF HOMELAND SECURITY CENTER OF EXCELLENCE




Viewpoint



Lighting



**Intra-Class
Differences**



Background



Deformation




Viewpoint



Lighting




Intra-Class Differences



Background



Deformation



Occlusion

Recap: KNN

[illegible]

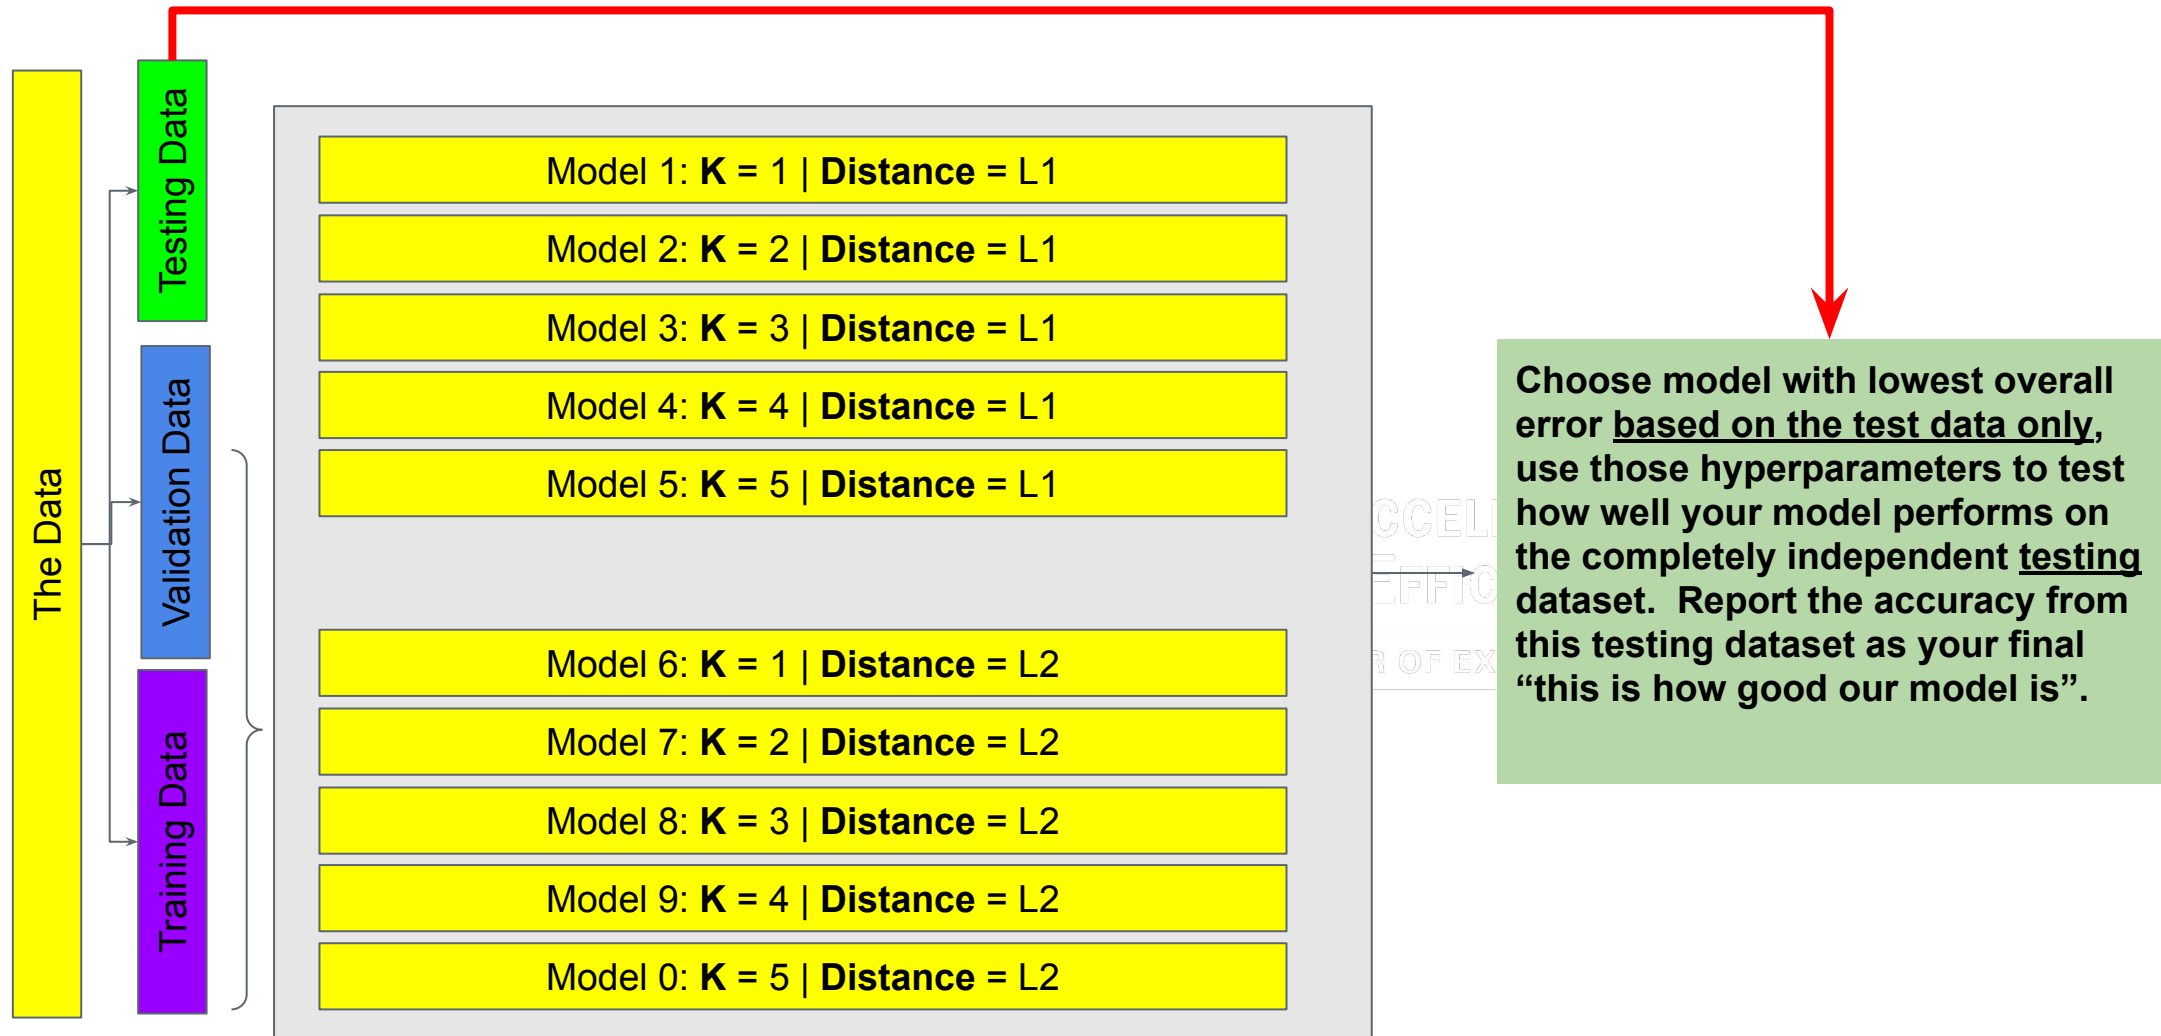
T from Training Data

[illegible]

T we want to Recognize

Sum of Absolute Difference: 10

[illegible]



Building Blocks of Neural Nets: Linear Classification

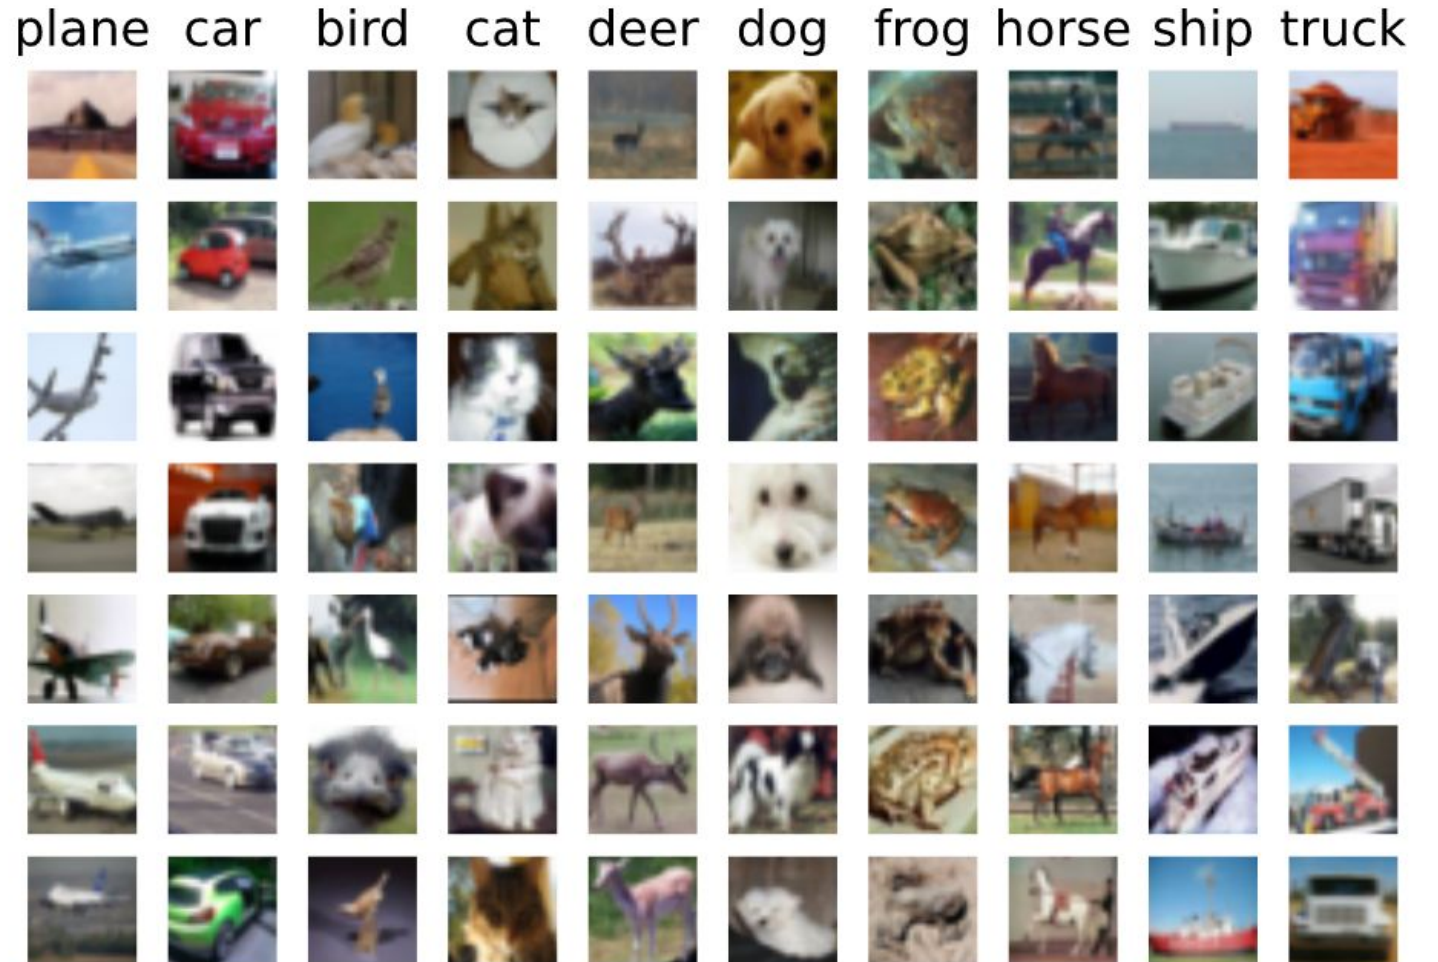
- Parametric vs. Non Parametric
- Interpreting Linear Classifiers
- Limitations of Linear Classifiers
- Segway into Loss Functions

CIFAR10 Dataset

(random examples generated from lab 1 code -->)

Goal: Given a new image, identify the correct class.

KNN approach: Record all of the images, and when a new image comes compare it to all images and select the most similar. Classify accordingly.





`nn.predict(image)`

	Probability
Bird	0.2
Dog	0.1
...	...
Cat	0.15
Plane	0.19

Parameters
(generally
referred to as
Weights)



`nn.predict(image, W)`

Probability	
Bird	0.2
Dog	0.1
...	...
Cat	0.15
Plane	0.19





def predict(image, W):
W*image

nn.predict(image, **W**)

	Probability
Bird	0.2
Dog	0.1
...	...
Cat	0.15
Plane	0.19

CIFAR10 Bird Example




$32 * 32 =$
1024 Pixels

32 Rows of Pixels

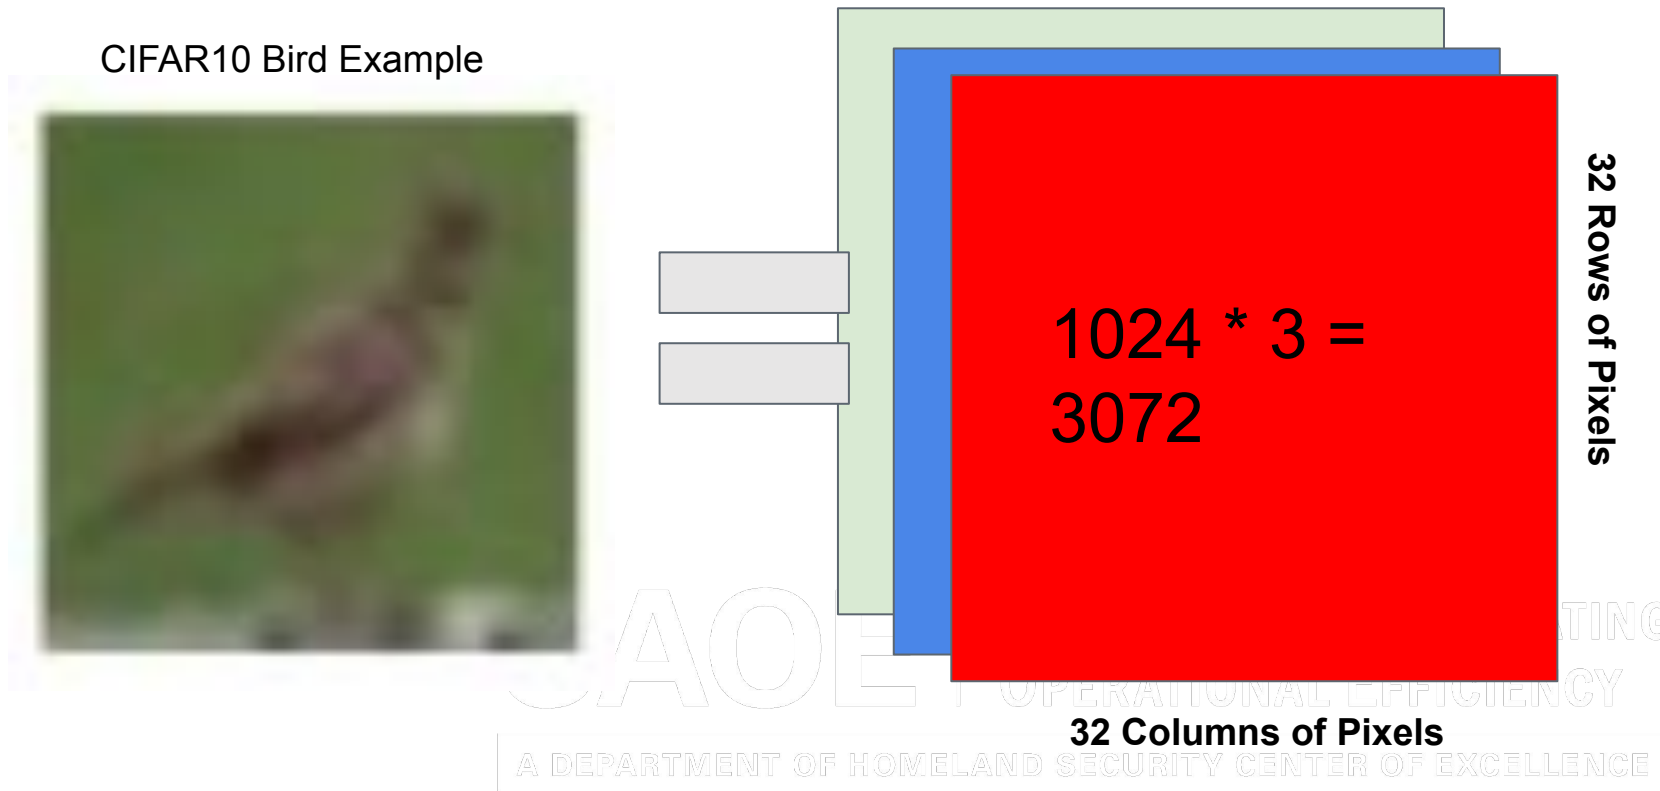
32 Columns of Pixels

CAOE CENTER FOR ADVANCING
OPERATIONAL EFFICIENCY
A DEPARTMENT OF HOMELAND SECURITY CENTER OF EXCELLENCE

CIFAR10 Bird Example




CIFAR10 Bird Example



```
def predict(image, w):  
    w*image
```

image: A vector of length 3072 - [0,12,3,2, 392] - where each value represents a pixel in one of the three color bands.


CIFAR10 Bird Example



```
def predict(image, W):  
     $W * image$ 
```

W: A 10x3072 matrix, with each of ten “columns” indicating the value to multiply by each pixel to generate a probability.


CIFAR10 Bird Example




```
def predict(image, W):  
    W*image
```


W*image: A 10 x 1 matrix in which each value is the probability of class inclusion.

CIFAR10 Bird Example






```
def predict(image, w):  
    w*image
```

```
def predict(image, W):
    W*image
```




0.2	-0.5	0.1	2.0	Cat
1.5	1.3	2.1	0.0	Bird
0	0.25	0.2	-0.3	Plane

$$\text{Cat Score} = (56 * 0.2) + (231 * -0.5) + (24 * 0.1) + (2 * 2.0) = -97.9$$

```
def predict(image, W):
```


$W * \text{image}$



0.2	-0.5	0.1	2.0	Cat
1.5	1.3	2.1	0.0	Bird
0	0.25	0.2	-0.3	Plane

$$\text{Cat Score} = (56 * 0.2) + (231 * -0.5) + (24 * 0.1) + (2 * 2.0) = -97.9$$

```
def predict(image, W):
    W*image
```



0.2	-0.5	0.1	2.0	Cat
1.5	1.3	2.1	0.0	Bird
0	0.25	0.2	-0.3	Plane


Cat Score = -97.9

Bird Score = 434.7

Plane Score = 63.15

GAOE | CENTER FOR ACCELERATING
OPERATIONAL EFFICIENCY
A DEPARTMENT OF HOMELAND SECURITY CENTER OF EXCELLENCE

Image (Matrix) to Vector



56
231
24
2

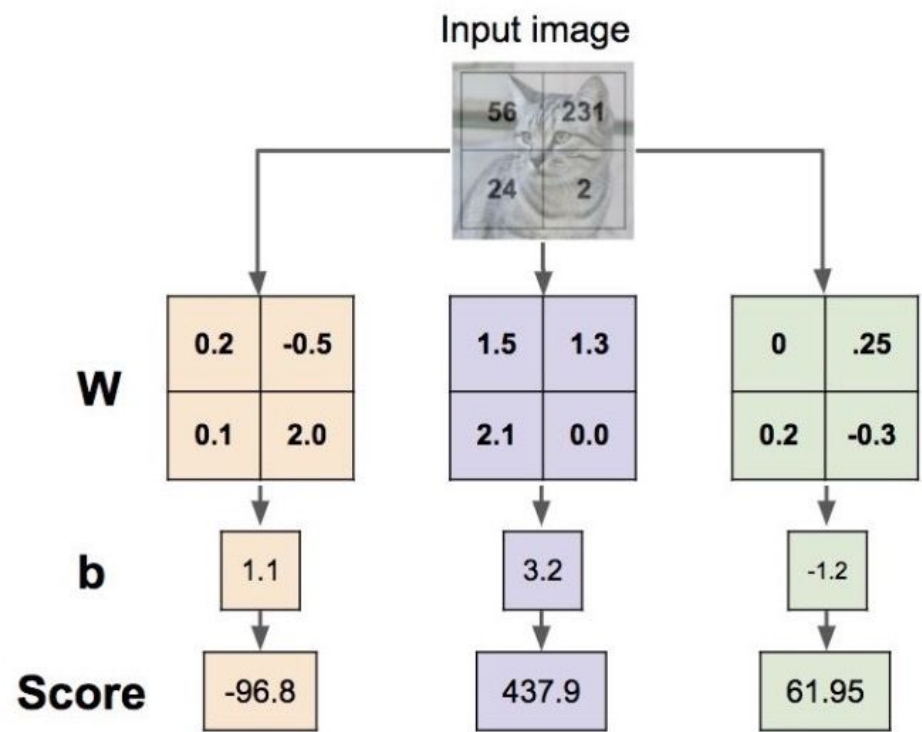

Weights Vector to Image

0.2	-0.5	0.1	2.0
-----	------	-----	-----

Cat →



airplane
 automobile
 bird
 cat
 deer
 dog
 frog
 horse
 ship
 truck



plane

car

bird

cat

deer

dog

frog


horse

ship

truck



CIFAR10 Bird Example



Loss Function


A single score that quantifies how bad a classification is.



Cat Score = -97.9

Bird Score = 3.5

Plane Score = 63.15





32 Rows of Pixels

32 Columns of Pixels

Optimization Strategy

**Finding the Weights
that minimize the loss
function.**



Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



$$f(\text{image}, W) = \text{scores}$$

CAOE | CENTER FOR ACCELERATING
OPERATIONAL EFFICIENCY
A DEPARTMENT OF HOMELAND SECURITY CENTER OF EXCELLENCE

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1




$$N=3 \{ (x_i, y_i) \}$$

3 images (indexed i=1, i=2, i=3).
Each image has image data (xi)
and a label (yi).

For example:


x1 =



y1 = "Cat"

f(image, W) = scores

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1





$$\text{Total Loss} = \frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i)$$

where **N** is the total number of images (i.e., 3), **i** is a unique index for each image, **x_i** is the image itself, **y_i** is the image label, **Loss_i** is the loss for that image, and **W** is the weights being tested.

f(image, W) = scores

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1




$$\text{Total Loss} = \frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i)$$

where **N** is the total number of images (i.e., 3), **i** is a unique index for each image, **x_i** is the image itself, **y_i** is the image label, **Loss_i** is the loss for that image, and **W** is the weights being tested.

f(image, W) = scores

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1




$$\text{Total Loss} = \frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i)$$

where **N** is the total number of images (i.e., 3), **i** is a unique index for each image, **x_i** is the image itself, **y_i** is the image label, **Loss_i** is the loss for that image, and **W** is the weights being tested.

f(image, W) = scores

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1




$$\text{Total Loss} = \frac{1}{N} \sum_i^N \boxed{Loss_i(f(x_i, W), y_i)}$$

where **N** is the total number of images (i.e., 3), **i** is a unique index for each image, **x_i** is the image itself, **y_i** is the image label, **Loss_i** is the loss for that image, and **W** is the weights being tested.

f(image, W) = scores

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1





$$\text{Total Loss} = \frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i)$$

where **N** is the total number of images (i.e., 3), **i** is a unique index for each image, **x_i** is the image itself, **y_i** is the image label, **Loss_i** is the loss for that image, and **W** is the weights being tested.

f(image, W) = scores

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



J is the total number of classes, represented by index j . In the current example, $j=1$ would be “Cat”, $j=2$ would be “Car”, etc.


s is the score for a given category. For the first image (the Cat), s_1 would be 3.2, s_2 would be 5.1, and s_3 would be -1.7.

Epsilon (ϵ) is a tolerance term, essentially defining how sure the algorithm needs to be about a class before we call it right.

Multiclass SVM Loss

$$\sum_{j \neq y_i}^J \max(0, s_j - s_{y_i} + \epsilon)$$

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



J is the total number of classes, represented by index j . In the current example, $j=1$ would be “Cat”, $j=2$ would be “Car”, etc.

s is the score for a given category. For the first image (the Cat), s_1 would be 3.2, s_2 would be 5.1, and s_3 would be -1.7.


Epsilon (ϵ) is a tolerance term, essentially defining how sure the algorithm needs to be about a class before we call it right.

$$\sum_{j \neq y_i}^J$$

$$\max(0, s_j - s_{y_i} + \epsilon)$$

Multiclass SVM Loss

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



J is the total number of classes, represented by index j . In the current example, $j=1$ would be “Cat”, $j=2$ would be “Car”, etc.


s is the score for a given category. For the first image (the Cat), s_1 would be 3.2, s_2 would be 5.1, and s_3 would be -1.7.

Epsilon (ϵ) is a tolerance term, essentially defining how sure the algorithm needs to be about a class before we call it right.

Multiclass SVM Loss

$$\sum_{j \neq y_i}^J \max(0, s_j - s_{y_i} + \epsilon)$$

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



J is the total number of classes, represented by index j . In the current example, $j=1$ would be “Cat”, $j=2$ would be “Car”, etc.



s is the score for a given category. For the first image (the Cat), s_1 would be 3.2, s_2 would be 5.1, and s_3 would be -1.7.

Epsilon (ϵ) is a tolerance term, essentially defining how sure the algorithm needs to be about a class before we call it right.

$$\sum_{j \neq y_i}^J \max(0, s_j - s_{y_i} + \epsilon)$$


Multiclass SVM Loss

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1




Epsilon (ϵ) is a tolerance term, essentially defining how sure the algorithm needs to be about a class before we call it right.

$$\sum_{j \neq y_i}^J \max(0, s_j - s_{y_i} + \epsilon)$$



Center for Accelerating
Operational Efficiency
OF HOMELAND SECURITY CENTER OF EXCELLENCE

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



If we set Epsilon = 1

Image X_1 (Cat) Loss:



$$\max(0, 5.1 - 3.2 + 1) = \max(0, 2.9) = 2.9$$

$$\sum_{j \neq y_i}^J \max(0, s_j - s_{y_i} + \varepsilon)$$

Multiclass SVM Loss

CAOE | CENTER FOR ACCELERATING
OPERATIONAL EFFICIENCY
A DEPARTMENT OF HOMELAND SECURITY CENTER OF EXCELLENCE

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



If we set Epsilon = 1

Image X_1 (Cat) Loss:

Car

$\max(0, 5.1 - 3.2 + 1) =$
 $\max(0, 2.9) =$
2.9



Frog

$\max(0, -1.7 - 3.2 + 1) =$
 $\max(0, -3.9) =$
0

$$\sum_{j \neq y_i}^J \max(0, s_j - s_{y_i} + \varepsilon)$$

Multiclass SVM Loss

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



If we set Epsilon = 1

Image X_2 (Car) Loss:

Cat

$$\max(0, 1.3 - 4.9 + 1) = \max(0, -2.6) = 0$$

Frog



$$\max(0, 2.0 - 4.9 + 1) = \max(0, -1.9) = 0$$

$$\sum_{j \neq y_i}^J \max(0, s_j - s_{y_i} + \varepsilon)$$

Multiclass SVM Loss

CAOE | CENTER FOR ACCELERATING
OPERATIONAL EFFICIENCY
A DEPARTMENT OF HOMELAND SECURITY CENTER OF EXCELLENCE

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



If we set Epsilon = 1

Image X_2 (Car) Loss:

Cat

$$\max(0, 2.2 - (-3.1) + 1) = \max(0, 6.3) = 6.3$$

Car



$$\max(0, 2.5 - (-3.1) + 1) = \max(0, -6.6) = 6.6$$

$$\sum_{j \neq y_i}^J \max(0, s_j - s_{y_i} + \varepsilon)$$

Multiclass SVM Loss

CAOE | CENTER FOR ACCELERATING OPERATIONAL EFFICIENCY
A DEPARTMENT OF HOMELAND SECURITY CENTER OF EXCELLENCE

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1





$$\text{Total Loss} = \frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i)$$

$$\sum_{j \neq y_i}^J \max(0, s_j - s_{y_i} + \varepsilon)$$

CAOE
CENTER FOR ACCELERATING
OPERATIONAL EFFICIENCY
A DEPARTMENT OF HOMELAND SECURITY CENTER OF EXCELLENCE

Loss	2.9	0	12.9
Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1





$$\text{Total Loss} = \frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i)$$

$$\sum_{j \neq y_i}^J \max(0, s_j - s_{y_i} + \epsilon)$$

$$(2.9 + 0 + 12.9) / 3 = \sim 5.27$$

Loss	2.9	0	12.9
Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



Wrap Up

- Parametric Models
- Linear Classifier
 - Solving
 - Visualizing
- Loss Functions
 - Multiclass SVM Loss

