

---

# DATA 442: Neural Networks & Deep Learning

Dan Runfola – [danr@wm.edu](mailto:danr@wm.edu)

[icss.wm.edu/data442/](https://icss.wm.edu/data442/)






def predict(image, W):  
W\*image

nn.predict(image, **W**)

	Probability
Bird	0.2
Dog	0.1
...	...
Cat	0.15
Plane	0.19



0.2	-0.5	0.1	2.0	Cat
1.5	1.3	2.1	0.0	Bird
0	0.25	0.2	-0.3	Plane

$$\text{Cat Score} = (56 * 0.2) + (231 * -0.5) + (24 * 0.1) + (2 * 2.0) = -97.9$$

```
def predict(image, W):
    W*image
```

Cat Score = -97.9

Bird Score = 434.7

Plane Score = 63.15




$$\text{Total Loss} = \frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i)$$

where **N** is the total number of images (i.e., 3), **i** is a unique index for each image, **x<sub>i</sub>** is the image itself, **y<sub>i</sub>** is the image label, **Loss<sub>i</sub>** is the loss for that image, and **W** is the weights being tested.

**f(image, W) = scores**

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



**J** is the total number of classes, represented by index  $j$ . In the current example,  $j=1$  would be “Cat”,  $j=2$  would be “Car”, etc.


**s** is the score for a given category. For the first image (the Cat),  $s_1$  would be 3.2,  $s_2$  would be 5.1, and  $s_3$  would be -1.7.

Epsilon ( $\epsilon$ ) is a tolerance term, essentially defining how sure the algorithm needs to be about a class before we call it right.

## Multiclass SVM Loss

$$\sum_{j \neq y_i}^J \max(0, s_j - s_{y_i} + \epsilon)$$

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



# Multinomial Logistic Regression - Softmax

Probability the picture is a


**Cat:** 34%

**Car:** 51%

**Frog:** 15%

**Softmax Function**

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1




# Multinomial Logistic Regression - Softmax

**Assumption:** These are really probabilities, just unnormalized!

**Specific Assumption:** These are unnormalized log probabilities for each class.

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1







$$P(Y = k | X = X_i)$$

**Assumption:** These are really probabilities, just unnormalized!

**Specific Assumption:** These are unnormalized log probabilities for each class.

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1





# Multinomial Logistic Regression - Softmax

$$P(Y = k | X = X_i) =$$


$$\frac{e^{s_k}}{\sum_{j=1}^J e^{s_j}}$$

This is the softmax function for class k.

**Assumption:** These are really probabilities, just unnormalized!

**Specific Assumption:** These are unnormalized log probabilities for each class.

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1





# Multinomial Logistic Regression - Softmax

$$\frac{e^s_k}{\sum_{j=1}^J e^s_j}$$

In a perfect world for this example, the above function would result in 1 for cat, and 0 for both car and frog.

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1




# Multinomial Logistic Regression - Softmax

$$Loss_i = -1 * \frac{e^s_k}{\sum_{j=1}^J e^s_j}$$

$$\frac{e^s_k}{\sum_{j=1}^J e^s_j}$$


Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



$$Loss_i = -1 * \log\left(\frac{e^s_k}{\sum_{j=1}^J e^s_j}\right)$$

$$\frac{e^s_k}{\sum_{j=1}^J e^s_j}$$

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1




# Multinomial Logistic Regression - Softmax

$$L_i = -\log\left(\frac{e^s_k}{\sum_{j=1}^J e^s_j}\right)$$

CAO Center for Accelerating  
EFFICIENCY  
A DEPARTMENT OF H  
OF EXCELLENCE

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1





# Multinomial Logistic Regression - Softmax

Class	Score	e^s
Cat	3.2	24.5
Car	5.1	164.0
Frog	-1.7	0.18

$$L_i = -log(\frac{e^s_k}{\sum_{j=1}^J e^s_j})$$

CAO Center for Accelerating  
EFFICIENCY  
A DEPARTMENT OF H  
OF EXCELLENCE

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1





# Multinomial Logistic Regression - Softmax


Class	Score	$e^s$	$e^s/188.68$
Cat	3.2	24.5	0.13
Car	5.1	164.0	0.87
Frog	-1.7	0.18	0.00

188.68

$$L_i = -\log\left(\frac{e^s_k}{\sum_{j=1}^J e^s_j}\right)$$

CAO Center for Accelerating  
EFFICIENCY  
A DEPARTMENT OF H  
OF EXCELLENCE

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



$$-\log(0.13) = 0.89$$

## Multinomial Logistic Regression - Softmax


Class	Score	$e^s$	$e^s/188.68$
Cat	3.2	24.5	0.13
Car	5.1	164.0	0.87
Frog	-1.7	0.18	0.00

188.68

$$L_i = -\log\left(\frac{e^s_k}{\sum_{j=1}^J e^s_j}\right)$$

CAO Center for Accelerating  
EFFICIENCY  
A DEPARTMENT OF H  
OF EXCELLENCE

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1





# Multinomial Logistic Regression - Softmax

$$L_i = -\log\left(\frac{e^s_k}{\sum_{j=1}^J e^s_j}\right)$$

CAOE  
OPERATIONAL EFFICIENCY  
A DEPARTMENT OF HOME AND SECURITY CENTER OF EXCELLENCE

Loss_i	0.89	0.034	2.67
Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1






## Multiclass SVM Loss


## Multinomial Logistic Regression - Softmax

$$\sum_{j \neq y_i}^J \max(0, s_j - s_{y_i} + \varepsilon) \quad L_i = -\log\left(\frac{e^{s_k}}{\sum_{j=1}^J e^{s_j}}\right)$$


<b>SVM</b>	<b>2.9</b>	<b>0</b>	<b>12.9</b>
<b>Softmax</b>	<b>0.89</b>	<b>0.034</b>	<b>2.67</b>
Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1














$$\frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i)$$

**Total Loss  
(Data Loss)**

CAOET CENTER FOR ACCELERATING  
OPERATIONAL EFFICIENCY

A DEPARTMENT OF HOMELAND SECURITY CENTER OF EXCELLENCE

$$\underbrace{\frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i)}_{\text{Data Loss}} + \underbrace{\lambda R(W)}_{\text{Regularization Loss}}$$

Data Loss

Regularization Loss



$$\underbrace{\frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i)}_{\text{Data Loss}} + \underbrace{\lambda R(W)}_{\text{Regularization Loss}}$$

Data Loss

Regularization Loss

CAOE | CENTER FOR ACCELERATING  
OPERATIONAL EFFICIENCY  
A DEPARTMENT OF HOMELAND SECURITY CENTER OF EXCELLENCE

$$\frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i) + \lambda R(W)$$

L2 Regularization

$$R(W) =$$

$$\sum_{k=1}^K w_k^2$$



$$\frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i) + \lambda R(W)$$

L1 Regularization


$$R(W) = \sum_{k=1}^K |W_k|$$

$$\frac{1}{N} \sum_i^N \text{Loss}_i(f(x_i, W), y_i) + \lambda R(W)$$

Elastic Net - Combination of L1 and L2

Max Norm Regularization, Batch Normalization, Stochastic Depths, Dropout Networks, ... many more.

# Summary

$$\prod_{i=1}^{N=3} \{(x_i, y_i)\}$$




# Summary

$$\prod_{i=1}^N \{ (x_i, y_i) \}$$



```
def predict(image, W):  
    return(W*image)
```



# Summary


$$\sum_{i=1}^{N=3} \{(x_i, y_i)\}$$



```
def predict(image, W):  
    return(W*image)
```



Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



# Summary


$$\sum_{i=1}^{N=3} \{(x_i, y_i)\}$$



```
def predict(image, W):  
    return(W*image)
```



Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



$$\rightarrow L_i = -\log\left(\frac{e_k^s}{\sum_{j=1}^J e_j^s}\right)$$



# Summary


Total Loss=

$$\sum_{i=1}^N \{(x_i, y_i)\}$$

$$\frac{1}{N} \sum_i^N Loss_i(f(x_i, W), y_i)$$

def predict(image, W):  
 return(W\*image)

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



$$L_i = -\log\left(\frac{e_k^s}{\sum_{j=1}^J e_j^s}\right)$$

# Summary

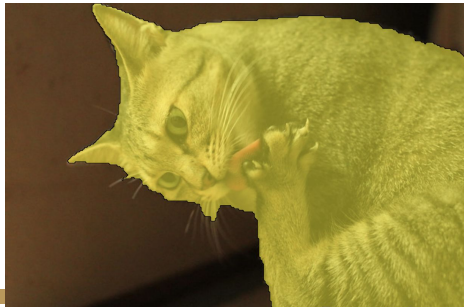
Total Loss=

$$\sum_{i=1}^N \{(x_i, y_i)\}$$

$$\frac{1}{N} \sum_i^N Loss_i(f(x_i, W), y_i) + \lambda R(W)$$

def predict(image, W):  
 return(W\*image)

Cat	3.2	1.3	2.2
Car	5.1	4.9	2.5
Frog	-1.7	2.0	-3.1



$$L_i = -\log\left(\frac{e_k^s}{\sum_{j=1}^J e_j^s}\right)$$